

# Support Vector Machines

Shan-Hung Wu  
*shwu@cs.nthu.edu.tw*

Department of Computer Science,  
National Tsing Hua University, Taiwan

NetDB-ML, Spring 2014

# Outline

- 1 Sparse Kernel Machines
- 2 Support Vector Regression
  - Primal Objective
- 3 Support Vector Classification
  - Primal Objective
  - Dual Problem
  - Sequential Minimal Optimization
  - Termination Criteria
  - Directional Search
  - Working Set Selection
  - Kernel Caching
- 4 Structural Risk Minimization

# Outline

- 1 Sparse Kernel Machines
- 2 Support Vector Regression
  - Primal Objective
- 3 Support Vector Classification
  - Primal Objective
  - Dual Problem
  - Sequential Minimal Optimization
  - Termination Criteria
  - Directional Search
  - Working Set Selection
  - Kernel Caching
- 4 Structural Risk Minimization

# Why Sparse Models?

- We have seen *dense* kernel machines 一般的Kernel Machine是dense
  - $h = \sum_{t=1}^N c_t k(\mathbf{x}^{(t)}, \cdot) = \sum_{t=1}^N c_t \Phi(\mathbf{x}^{(t)})$  where  $c_t \neq 0$  for most  $t$
- To make a prediction  $h(\mathbf{x}') = \sum_{t=1}^N c_t k(\mathbf{x}^{(t)}, \mathbf{x}')$ , we need to store *all* examples
- May be infeasible in practice due to
  - Big dataset (large  $N$ )
  - Time limit
  - Space limit
- Solution: make the kernel machines *sparse*;
  - I.e., make  $c_t \neq 0$  for only a small fraction of the examples called *support vectors*
- How?

# Creating Sparsity

如果 $r^1$ 恆等於(任何 $g$ ) $g(x^1)$ ，並且constraints也都滿足，  
那麼 $c^1=0$ ，因為 $x^1$ 並不影響 $h$

- Recall that a kernel machine  $h = \sum_{t=1}^N c_t k(\mathbf{x}^{(t)}, \cdot)$  is the solution to

$$\begin{aligned} \arg \min_{g \in \mathcal{H}} L((\mathbf{x}^{(1)}, r^{(1)}, g(\mathbf{x}^{(1)})), \dots, (\mathbf{x}^{(N)}, r^{(N)}, g(\mathbf{x}^{(N)}))) + \Omega(\|g\|_{\mathcal{RKHCS}}) \\ \text{subject to } C_k((\mathbf{x}^{(1)}, r^{(1)}, g(\mathbf{x}^{(1)})), \dots, (\mathbf{x}^{(N)}, r^{(N)}, g(\mathbf{x}^{(N)}))) \leq 0, \forall k \end{aligned}$$

- In what situation, will we have  $c_1 = 0$ ? If for different  $g$ ,  $g(\mathbf{x}^{(1)})$  neither changes the value of  $L$  nor violates the constraints  $C_k(\dots) \leq 0$ , then minimizing the regularization term makes  $c_1 = 0$
- Idea:
  - Suppose  $L((\mathbf{x}^{(1)}, r^{(1)}, g(\mathbf{x}^{(1)})), \dots, (\mathbf{x}^{(N)}, r^{(N)}, g(\mathbf{x}^{(N)}))) = \sum_{t=1}^N l(\mathbf{x}^{(t)}, r^{(t)}, g(\mathbf{x}^{(t)}))$ , refine  $l$  such that  $l(\mathbf{x}^{(t)}, r^{(t)}, g(\mathbf{x}^{(t)})) = 0$  for most  $t$ 's no matter what the  $g$  is
  - Similarly, refine constraints (if any) such that they hold for most  $t$ 's no matter what the  $g$  is

# Outline

- 1 Sparse Kernel Machines
- 2 **Support Vector Regression**
  - Primal Objective
- 3 **Support Vector Classification**
  - Primal Objective
  - Dual Problem
  - Sequential Minimal Optimization
  - Termination Criteria
  - Directional Search
  - Working Set Selection
  - Kernel Caching
- 4 **Structural Risk Minimization**

# Support Vector Regression

$\epsilon$ : Hyper-parameter, 誤差容忍

$\text{error} < \epsilon \rightarrow \text{error} = 0$


- Idea: to create an  **$\epsilon$ -insensitive loss**: 太小的error就當作沒error

$$\arg \min_{\mathbf{w}, b, \xi^+, \xi^-} \sum_{t=1}^N (\xi_t^+ + \xi_t^-) + \lambda \|\mathbf{w}\|^2$$

subject to  $(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) - r^{(t)} \leq \epsilon + \xi_t^+$

$\xi_t^+, \xi_t^-$ : 離預測最遠的距離  $(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) - r^{(t)} \geq -\epsilon - \xi_t^-$

$$\xi_t^+ \geq 0, \quad \xi_t^- \geq 0, \quad \forall t = 1, \dots, N$$



- $\lambda$  and  $\epsilon$  are hyperparameters
- No matter what the  $\mathbf{w}$  and  $b$  are, most examples (those falling inside the “ $\epsilon$ -tube”) will have  $\xi_t^+ = 0$ ,  $\xi_t^- = 0$ , and the constraints always hold  
「誤差免除」距離以內
- So, if the representation theorem applies, the solution will be sparse
  - Examples falling outside the  $\epsilon$ -tube are support vectors
- Is the representation theorem applicable?  $\xi_t^+, \xi_t^-$ : 離預測太遠的資料會影響 $\mathbf{w}$ 的走向

# Applying Representer Theorem

- Note that  $\xi_t^+ \geq 0$  and  $(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) - r^{(t)} \leq \epsilon + \xi_t^+$  implies  
 $\xi_t^+ = \max(0, (\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) - (r^{(t)} + \epsilon))$  預測值偏差太大→正數，否則0
- Similarly,  $\xi_t^- = \max(0, (r^{(t)} - \epsilon) - (\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b))$  預測值偏差太小→正數，否則0
- We can eliminate  $\xi^+$  and  $\xi^-$  by rewriting the objective as

$$\arg \min_{\mathbf{w}, b} \sum_{t=1}^N (\max(0, (\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) - (r^{(t)} + \epsilon)) + \max(0, (r^{(t)} - \epsilon) - (\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b))) + \lambda \|\mathbf{w}\|^2$$

- Convex!** ( $\max(a, b)$  is convex if  $a$  and  $b$  are convex)
- Shape of the loss function  $l$ ? 線性：既 Convex 也 Concave  
Maximize a convex: convex
- The semiparametric representation theorem applies here:

$\|\mathbf{w}\|^2$ : convex

- $h = \sum_{t=1}^N c_t k(\mathbf{x}^{(t)}, \cdot) + b$
- Feature-dimension-independent problem:

$$\arg \min_{\mathbf{c} \in \mathbb{R}^N, b} \sum_{t=1}^N (\max(0, (\mathbf{K}_{t,\cdot} \mathbf{c} + b) - (r^{(t)} + \epsilon)) + \max(0, (r^{(t)} - \epsilon) - (\mathbf{K}_{t,\cdot} \mathbf{c} + b))) + \lambda \mathbf{c}^\top \mathbf{K} \mathbf{c}$$

- $\mathbf{K}_{t,\cdot}$  is the  $t$ -th row of  $\mathbf{K}$

Regularized，  
並且使用  $\epsilon$ -tube，  
讓越少的  $\mathbf{x}$  參與決策  
(成為 Support Vector)  
就可以讓越多  $K_t$  是 0  
從而產生稀疏性



# Outline

- 1 Sparse Kernel Machines
- 2 Support Vector Regression
  - Primal Objective
- 3 Support Vector Classification**
  - Primal Objective
  - Dual Problem
  - Sequential Minimal Optimization
  - Termination Criteria
  - Directional Search
  - Working Set Selection
  - Kernel Caching
- 4 Structural Risk Minimization

# Outline

- 1 Sparse Kernel Machines
- 2 Support Vector Regression
  - Primal Objective
- 3 **Support Vector Classification**
  - **Primal Objective**
  - Dual Problem
  - Sequential Minimal Optimization
  - Termination Criteria
  - Directional Search
  - Working Set Selection
  - Kernel Caching
- 4 Structural Risk Minimization

# Support Vector Classification (1)

- The idea of  $\epsilon$ -insensitive loss still results in dense solution when applied to the classification problem. Why?
  - Model is linear in feature space, where
  - examples close to or far away from the decision boundary (i.e.,  $\{\mathbf{x} : h(\mathbf{x}) = 0\}$ , a hyperplane) will all be support vectors 與需求相反
- Holds even if we adopt a non-linear kernel
  - Non-linear kernel “re-positions” instances (by creating feature space)
  - Makes instances more separable
  - But does not change the label values  $+1$ 's and  $-1$ 's 會這樣做
- Any new idea?

# Support Vector Classification (2)

- We can borrow the idea of the perceptron to remove support vectors far from the decision boundary:

$$\begin{aligned} & \arg \min_{\mathbf{w}, b, \xi} \sum_{t=1}^N \xi_t + \lambda \|\mathbf{w}\|^2 \\ & \text{subject to } r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) \geq -\xi_t \text{ and } \xi_t \geq 0, \forall t = 1, \dots, N \end{aligned}$$

- Examples not on the decision boundary nor misclassified will have  $\xi_t = 0$  and the constraints hold 被分得「太對」的，不會是Support Vector
- Only examples on the decision boundary and outliers will be support vectors 差點被分錯，或是已經被分錯的，才會是Support Vector


# Support Vector Classification (3)

$$\begin{aligned} & \arg \min_{\mathbf{w}, b, \xi} \sum_{t=1}^N \xi_t + \lambda \|\mathbf{w}\|^2 \\ & \text{subject to } r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) \geq -\xi_t \text{ and } \xi_t \geq 0, \forall t = 1, \dots, N \end{aligned}$$

- With a lifting/kernel function, instances are more separable
- The above objective has problems in dealing with a perfectly separable dataset:

- ① Multiple hyperplanes can be equally good. Which one is the best?
- ② The hyperplane could be arbitrary flat, causing numerical problems

因為要盡可能讓 $\|\mathbf{w}\|$ 小一點，所以這條線會比較平，但其實不應該這麼平



3條線都不會產生emp error，但左右2條只要有東西落在靠中間，很容易就分錯了


- Solution?

# Support Vector Classification (4)

- Refined objective:

$$\arg \min_{\mathbf{w}, b, \xi} \sum_{t=1}^N \xi_t + \lambda \|\mathbf{w}\|^2 \quad \text{找盡可能smooth的}\mathbf{w}$$

subject to  $r^{(t)}((\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) - r^{(t)}) \geq -\xi_t$  and  $\xi_t \geq 0, \forall t = 1, \dots, N$   
並且 讓正examples和負examples離預測零點有一定的、相同的距離  
即、把線放在最負的正examples和最正的負examples的正中間



- Always find the hyperplane that sits “in the middle” between positive and negative examples
- Not arbitrarily flat anymore: the flatter the hyperplane, the more the slacks  
Slack variables, 鬆弛參數,  $\xi^{(t)}$  是用來衡量標記錯誤的樣本、其錯誤的程度

# Large Margin Perspective

- Form geometric point of view, minimizing the term  $\|\mathbf{w}\|^2$  in the problem:


C是Regularize用的  
Hyper-Parameter

$$\arg \min_{\mathbf{w}, b, \xi} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{t=1}^N \xi_t$$

subject to  $r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) \geq 1 - \xi_t$  and  $\xi_t \geq 0, \forall t = 1, \dots, N$

amounts to finding a hyperplane that *maximizes the margin*

意味著



- The distance between  $\{\mathbf{x} : \mathbf{w}^\top \mathbf{x} + b = 1\}$  and  $\{\mathbf{x} : \mathbf{w}^\top \mathbf{x} + b = -1\}$  is  $\frac{2}{\|\mathbf{w}\|}$  [Homework]
- Examples touching or falling inside the margin are support vectors
- The distance between  $\{\mathbf{x} : \mathbf{w}^\top \mathbf{x} + b = a\}$  and  $\{\mathbf{x} : \mathbf{w}^\top \mathbf{x} + b = -a\}$  is  $\frac{2a}{\|\mathbf{w}\|}$ . Why pick  $a = 1$ ? Changing  $a$  does *not* change the hyperplane we find 不管a怎麼選，重點是要minimize  $\|\mathbf{w}\|$

# Applying Representer Theorem

- Again,  $r^{(t)}((\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) - r^{(t)}) \geq -\xi_t$  and  $\xi_t \geq 0$  implies  
 $\xi_t = \max(0, -r^{(t)}((\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) - r^{(t)}))$  跑對邊,  $\xi_t=0$  跑錯邊,  $\xi_t$ =錯得多離譜
- Also,  $r^{(t)}((\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) - r^{(t)})$  can be written as  
 $r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) - 1$   $r^t r^t = 1$
- Objective with  $\xi_t$  eliminated:

$$\arg \min_{\mathbf{w}, b} \sum_{t=1}^N \max(0, 1 - r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b)) + \lambda \|\mathbf{w}\|^2$$

- Semiparametric representation theorem applies:  
 $h = \sum_{t=1}^N c_t k(\mathbf{x}^{(t)}, \cdot) + b$
- Feature-dimension-independent objective:

$$\arg \min_{\mathbf{c} \in \mathbb{R}^N, b} \sum_{t=1}^N \max(0, 1 - r^{(t)}(\mathbf{K}_{t, \cdot} \mathbf{c} + b)) + \lambda \mathbf{c}^\top \mathbf{K} \mathbf{c}$$

- Still **convex**



# Outline

- 1 Sparse Kernel Machines
- 2 Support Vector Regression
  - Primal Objective
- 3 Support Vector Classification**
  - Primal Objective
  - Dual Problem**
  - Sequential Minimal Optimization
  - Termination Criteria
  - Directional Search
  - Working Set Selection
  - Kernel Caching
- 4 Structural Risk Minimization

# Dual Problem

- Primal problem:

最平滑(Margin最大) · 分錯的錯誤(slack)總和最低

$$\begin{aligned} & \arg \min_{\mathbf{w}, b, \xi} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{t=1}^N \xi_t \\ & \text{subject to } r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) \geq 1 - \xi_t \text{ and } \xi_t \geq 0, \forall t = 1, \dots, N \\ & \text{對於錯誤的容忍} \end{aligned}$$

- Dual problem:

先找一條線分開  
再把margin掰大

$$\begin{aligned} & \arg \max_{\alpha, \beta} \inf_{\mathbf{w}, b, \xi} \mathcal{L}(\mathbf{w}, b, \xi, \alpha, \beta) \\ & \text{subject to } \alpha \geq 0, \beta \geq 0 \end{aligned}$$

where  $\mathcal{L}(\mathbf{w}, b, \xi, \alpha, \beta) =$  Constraints with Lagrange Multipliers

$$\frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{t=1}^N \xi_t + \sum_{t=1}^N \alpha_t (1 - r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) - \xi_t) + \sum_{t=1}^N \beta_t (-\xi_t)$$

- Why dual? We can distinguish SVs **on** the margin  $\{\mathbf{x}^{(t)} : r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) = 1\}$  from those **inside** the margin  $\{\mathbf{x}^{(t)} : r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) < 1\}$
- How to solve it?

# Solving the Dual Problem

偏微分，導出最優化條件

- $dual(\alpha, \beta) = \inf_{\mathbf{w} \in \mathbb{R}^n} \mathcal{L}(\mathbf{w}, b, \xi, \alpha, \beta)$  is convex in terms of  $\mathbf{w}$ ,  $b$ , and  $\xi_t$
- $\frac{\partial \mathcal{L}}{\partial \mathbf{w}} = \mathbf{w} - \sum_{t=1}^N \alpha_t r^{(t)} \mathbf{x}^{(t)} = 0 \Rightarrow \mathbf{w} = \sum_{t=1}^N \alpha_t r^{(t)} \Phi(\mathbf{x}^{(t)})$
- $\frac{\partial \mathcal{L}}{\partial b} = \sum_{t=1}^N \alpha_t r^{(t)} = 0$  零和 (zero-sum)
- $\frac{\partial \mathcal{L}}{\partial \xi_t} = C - \alpha_t - \beta_t = 0 \Rightarrow \beta_t = C - \alpha_t$  對每一個Example而言，Slack的大小決定 $\alpha_t$ 和C的大小關係
  - Additionally, since  $\beta_t \geq 0$ , we have  $\alpha_t \leq C$

- $dual(\alpha, \beta; \mathbf{x}) = \begin{cases} \sum_{t=1}^N \alpha_t - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j r^{(i)} r^{(j)} \Phi(\mathbf{x}^{(i)})^\top \Phi(\mathbf{x}^{(j)}), & \text{if } \sum_{t=1}^N \alpha_t r^{(t)} = 0 \\ -\infty, & \text{otherwise} \end{cases}$  最優化條件  
就不會被外面的maximize選到  
 $\forall i, \alpha_i \leq C$
- Dual problem:

$$\begin{aligned} & \arg \max_{\alpha} \alpha^\top \mathbf{1} - \frac{1}{2} \alpha^\top \tilde{\mathbf{K}} \alpha \\ & \text{subject to } \mathbf{0} \leq \alpha \leq C \mathbf{1}, \mathbf{r}^\top \alpha = 0 \end{aligned}$$

- $\tilde{\mathbf{K}}_{i,j} = r^{(i)} r^{(j)} k(\mathbf{x}^{(i)}, \mathbf{x}^{(j)})$

# Primal vs. Dual

- Primal problem:


$$\begin{aligned} \arg \min_{\mathbf{w}, b, \xi} \text{primal}(\mathbf{w}, b, \xi) &:= \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_t \xi_t \\ \text{subject to } r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) &\geq 1 - \xi_t \text{ and } \xi_t \geq 0, \forall t = 1, \dots, N \end{aligned}$$

- Dual problem:

$$\begin{aligned} \arg \max_{\alpha} \text{dual}(\alpha) &:= \alpha^\top \mathbf{1} - \frac{1}{2} \alpha^\top \tilde{\mathbf{K}} \alpha = \sum_t \alpha_t - \frac{1}{2} \sum_{i,j} r^{(i)} \alpha_i r^{(j)} \alpha_j \mathbf{K}_{i,j} \\ \text{subject to } \mathbf{0} &\leq \alpha \leq C \mathbf{1}, r^\top \alpha = 0 \end{aligned}$$

$C$  是 Hyper-Parameter

- Why solving dual problem?
- Constrains become much simpler: only **box** constraint and **zero-sum** constraint
- Better understanding of the data



# Types of SVs

C是Hyper-Parameter

- By  $\mathbf{w} = \sum_{t=1}^N \alpha_t r^{(t)} \Phi(\mathbf{x}^{(t)})$ , examples whose  $\alpha_t > 0$  will be support vectors (SVs)
  - Non-SVs are those examples whose  $\alpha_t = 0$  margin外  $r^{(t)}\Phi(\mathbf{x}^{(t)}) > 1 : \alpha_t = 0$
  - **Free SVs**:  $0 < \alpha_t < C$  margin上  $r^{(t)}\Phi(\mathbf{x}^{(t)}) = 1 : \alpha_t = [0, C]$
  - **Bounded SVs**:  $\alpha_t = C$  margin内  $r^{(t)}\Phi(\mathbf{x}^{(t)}) < 1 : \alpha_t = C$
- Where do the corresponding examples  $\mathbf{x}^{(t)}$  appear in the graph of  $f(\cdot) = \mathbf{w}^\top \Phi(\cdot) + b$ ?

# KKT Conditions

- By KKT conditions, we have:

**定義** • Primal feasibility:  $r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) \geq 1 - \xi_t$  and  $\xi_t \geq 0, \forall t = 1, \dots, N$

**Dual** • Complementary slackness:  $\alpha_t(1 - r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) - \xi_t) = 0$  and  
**正確運作**  $\beta_t(-\xi_t) = 0$  無Slack:非SV,  $\alpha_t=0$  有Slack:是SV,  $\alpha_t=C$

- Non SVs whose  $\alpha_t = 0$ :  $r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) \geq 1$  **C是Hyper-Parameter**
  - $1 - r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) - \xi_t \leq 0$  margin外(非Support Vector)
  - Since  $\beta_t = C - \alpha_t \neq 0$ , we have  $\xi_t = 0$   $r^{(t)}\Phi(\mathbf{x}^{(t)}) > 1 : \alpha_t = 0$
- Free SVs** whose  $0 < \alpha_t < C$ :  $r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) = 1$ 
  - $1 - r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) - \xi_t = 0$  margin上
  - Since  $\beta_t = C - \alpha_t \neq 0$ , we have  $\xi_t = 0$   $r^{(t)}\Phi(\mathbf{x}^{(t)}) = 1 : \alpha_t = [0, C]$
- Bounded SVs** whose  $\alpha_t = C$ :  $r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) \leq 1$  (usually strick)
  - $1 - r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) - \xi_t = 0$  margin內
  - Since  $\beta_t = C - \alpha_t = 0$ , hence  $\xi_t \geq 0$   $r^{(t)}\Phi(\mathbf{x}^{(t)}) < 1 : \alpha_t = C$

差點被分錯邊(只要不夠對), 就有slack了(真的分錯也有)

# Solving $b$

- How to obtain  $b$ ?
- By complementary slackness in KKT conditions, we have  $\alpha_t(1 - r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) - \xi_t) = 0$  for any  $t$
- For a free SV  $\mathbf{x}^{(t)}$  (whose  $0 < \alpha_t < C$ , thus  $\xi_t = 0$ ), we have  $1 = r^{(t)}(\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) \Rightarrow b = r^{(t)} - \mathbf{w}^\top \Phi(\mathbf{x}^{(t)})$ 
  - Can be computed based on **any** free SV
- In practice, we usually take the average over **all** free SVs to avoid numeric error

# Outline

- 1 Sparse Kernel Machines
- 2 Support Vector Regression
  - Primal Objective
- 3 Support Vector Classification**
  - Primal Objective
  - Dual Problem
  - Sequential Minimal Optimization**
  - Termination Criteria
  - Directional Search
  - Working Set Selection
  - Kernel Caching
- 4 Structural Risk Minimization



# SVM Solvers

- SVM primal/dual objectives can be solved by standard convex or quadratic programming software
- In practice, dedicated solvers are developed, since
- Quadratic optimization packages were often designed to take advantage of sparsity in the quadratic part of the objective function
  - Unfortunately, the SVM kernel matrix is rarely sparse; sparsity occurs in the solution of the SVM problem
- The specification of a SVM problem rarely fits in memory. Kernel matrix coefficient must be cached or computed on the fly
  - Vast speedups are achieved by accessing the kernel matrix coefficients carefully
- Generic optimization packages sometimes make extra work to locate the optimum with high accuracy
  - The accuracy requirements of a learning problem are unusually low

# Dual Problem Revisited

- Primal problem:

$$\begin{aligned} \arg \min_{\mathbf{w}, b, \xi} \text{primal}(\mathbf{w}, b, \xi) &:= \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_t \xi_t \\ \text{subject to } r^{(t)} (\mathbf{w}^\top \Phi(\mathbf{x}^{(t)}) + b) &\geq 1 - \xi_t \text{ and } \xi_t \geq 0, \forall t = 1, \dots, N \end{aligned}$$

- Dual problem:

$$\begin{aligned} \arg \max_{\alpha} \text{dual}(\alpha) &:= \alpha^\top \mathbf{1} - \frac{1}{2} \alpha^\top \tilde{\mathbf{K}} \alpha = \sum_t \alpha_t - \frac{1}{2} \sum_{i,j} r^{(i)} \alpha_i r^{(j)} \alpha_j \mathbf{K}_{i,j} \\ \text{subject to } \mathbf{0} &\leq \alpha \leq C \mathbf{1}, \mathbf{r}^\top \alpha = 0 \end{aligned}$$

- **Box** constraint and **zero-sum** constraint

# Sequential Minimal Optimization

- A special case of *coordinate descent* (or *decomposition methods*) which picks only few coordinates to update during each iteration
- What's the minimal #coordinates for SVM dual? 2, due to the zero-sum constraint  $\mathbf{r}^\top \boldsymbol{\alpha} = 0$

Repeat until convergence {

1. Select some coordinate pair  $\alpha_i$  and  $\alpha_j$  to update next using some heuristic (e.g., to pick the two that allow us to make the biggest progress towards the global maximum); 選讓最佳化問題可以最有進展的兩個點

2. Re-optimize  $dual(\boldsymbol{\alpha})$  with respect to  $\alpha_i$  and  $\alpha_j$ , while holding all the other coordinates  $\alpha_k$ 's,  $k \neq i, j$ , fixed;

}

- As compared to the batch descent: more iterations, but each iteration can run very fast!

相較於Gradient decent，雖然iteration多，但是每個iteration快

# Outline

- 1 Sparse Kernel Machines
- 2 Support Vector Regression
  - Primal Objective
- 3 Support Vector Classification**
  - Primal Objective
  - Dual Problem
  - Sequential Minimal Optimization
  - Termination Criteria**
  - Directional Search
  - Working Set Selection
  - Kernel Caching
- 4 Structural Risk Minimization

# Optimality Conditions Revisited (1)

- It is sometimes convenient to rewrite the box constraint  $0 \leq \alpha_t \leq C$  as:

先不要管Free

Support Vectors

$$r^{(t)} \alpha_t \in [A^{(t)}, B^{(t)}] = \begin{cases} [0, C], & r^{(t)} = 1 \\ [-C, 0], & r^{(t)} = -1 \end{cases}$$

總之乘起來  
就是在  
0~±C之間

- Let  $\alpha^*$  be the solution to the dual problem
- Consider a pair of subscripts  $(i, j)$  such that  $r^{(i)} \alpha_i^* < B^{(i)}$  and  $A^{(j)} < r^{(j)} \alpha_j^*$
- Define  $\alpha^\epsilon$  such that  $\alpha_t^\epsilon = \alpha_t^* + \begin{cases} \epsilon r^{(t)}, & \text{if } t = i \\ -\epsilon r^{(t)}, & \text{if } t = j \\ 0, & \text{otherwise} \end{cases}$ , clearly  $\alpha^\epsilon$  is also a feasible solution when  $\epsilon$  is positive and sufficiently small

$I_{up}$ ：在負Margin上方

正Instance： $r^{(i)}\alpha_i < C$ ， $\alpha_i < C$ ， $r^{(i)}$ 在線上或線外

負Instance： $r^{(i)}\alpha_i < 0$ ， $\alpha_i > 0$ ， $r^{(i)}$ 在線上或線內

$I_{down}$ ：在正Margin下方


正Instance： $r^{(i)}\alpha_i > 0$ ， $\alpha_i > 0$ ， $r^{(i)}$ 在線上或線內

負Instance： $r^{(i)}\alpha_i > -C$ ， $\alpha_i < C$ ， $r^{(i)}$ 在線上或線外

By definition

$$I_{up} := \{ i \mid r^{(i)}\alpha_i^* < B^{(i)} \}$$

$$I_{down} := \{ j \mid A^{(j)} < r^{(j)}\alpha_j^* \}$$



# Optimality Conditions Revisited (2)

- Also, define the **gradient**  $g_t$  of  $dual$  at  $\alpha$  as 泰勒展開式的前2項

$$g_t := \frac{\partial dual(\alpha)}{\partial \alpha_t} = 1 - r^{(t)} \sum_i r^{(i)} \alpha_i K_{t,i} \quad [\text{Proof}]$$

- By the first order expansion, we have

$$dual(\alpha^\epsilon) - dual(\alpha^*) = \epsilon(r^{(i)} g_i^* - r^{(j)} g_j^*) + o(\epsilon)$$

- $r^{(i)} g_i^* - r^{(j)} g_j^* \leq 0$  as  $dual(\alpha^\epsilon) \leq dual(\alpha^*)$  因為 $\alpha^*$ 已是最佳解下的 $\alpha$

- Since this holds for all pairs  $(i, j)$  such that  $r^{(i)} \alpha_i^* < B^{(i)}$  and

$A^{(j)} < r^{(j)} \alpha_j^*$ , we have another **necessary optimality condition**:


$$\exists \rho \in \mathbb{R} \text{ such that } \max_{i \in I_{up}} r^{(i)} g_i^* \leq \rho \leq \min_{j \in I_{down}} r^{(j)} g_j^*$$

where  $I_{up} := \{i | r^{(i)} \alpha_i^* < B^{(i)}\}$  and  $I_{down} := \{j | A^{(j)} < r^{(j)} \alpha_j^*\}$


- $\rho$  is unique, due to the existence of free SVs that exist in both  $I_{up}$  and

$I_{down}$   $I_{up} \square\square\square\square\square\square I_{down} \square\square\square\square\square\square$  這兩個之間的差距

Before start



End



$$\exists \rho \in \mathbb{R} \text{ such that } \forall t, \begin{cases} r^{(t)} \alpha_t^* = B^{(t)}, & \text{if } r^{(t)} g_t^* > \rho \\ r^{(t)} \alpha_t^* = A^{(t)}, & \text{if } r^{(t)} g_t^* < \rho \end{cases}$$

$$\exists \rho \in \mathbb{R} \text{ such that } \forall t, \begin{cases} \alpha_t^* = C, & \text{if } g_t^* > r^{(t)} \rho \\ \alpha_t^* = 0, & \text{if } g_t^* < r^{(t)} \rho \end{cases}$$

If  $r^{(t)} g_t^* > \rho$  and  $r^{(t)} \alpha_t^* < B^{(t)}$ , then the instance  $t$  must be in  $I_{up}$  (consider instance  $x_1$ ), meaning  $\max_{i \in I_{up}} r^{(i)} g_i > \rho$  in Condition 1, a contradiction. Similarly, if  $r^{(t)} g_t^* < \rho$  and  $r^{(t)} \alpha_t^* > A^{(t)}$ , then the instance  $t$  must be in  $I_{down}$  (consider instance  $x_2$ ), leading to another contradiction  $\rho > \min_{j \in I_{down}} r^{(j)} g_j$  to Condition 1.



# Optimality Conditions Revisited (3)

- Actually the above condition is also *sufficient*
- By letting  $\mathbf{w}^* := \sum_{t=1}^N \alpha_t^* r^{(t)} \Phi(\mathbf{x}^{(t)})$ ,  $b^* := \rho$ , and  $\xi_t^* := \max\{0, g_t^* - r^{(t)} \rho\}$ , we can see that [Proof]
  - All primal constraints are satisfied, and
  - $\text{primal}(\mathbf{w}^*, b^*, \xi^*) - \text{dual}(\boldsymbol{\alpha}^*) = 0$ , so no duality gap

# Sequential Minimal Optimization

---

**Algorithm 6.2** SMO with Maximum Violating Pair working set selection

---

```
1:  $\forall k \in \{1 \dots n\} \quad \alpha_k \leftarrow 0$  ▷ Initial coefficients  
2:  $\forall k \in \{1 \dots n\} \quad g_k \leftarrow 1$  ▷ Initial gradient  
3: loop  
4:    $i \leftarrow \arg \max_i y_i g_i \quad \text{subject to } y_i \alpha_i < B_i$   
5:    $j \leftarrow \arg \min_j y_j g_j \quad \text{subject to } A_j < y_j \alpha_j$  ▷ Maximal violating pair  
6:   if  $y_i g_i \leq y_j g_j$  stop. ▷ Optimality criterion (11)  
7:    $\lambda \leftarrow \min \left\{ B_i - y_i \alpha_i, y_j \alpha_j - A_j, \frac{y_i g_i - y_j g_j}{K_{ii} + K_{jj} - 2K_{ij}} \right\}$  ▷ Direction search  
8:    $\forall k \in \{1 \dots n\} \quad g_k \leftarrow g_k - \lambda y_k K_{ik} + \lambda y_k K_{jk}$  ▷ Update gradient  
9:    $\alpha_i \leftarrow \alpha_i + y_i \lambda \quad \alpha_j \leftarrow \alpha_j - y_j \lambda$  ▷ Update coefficients  
10: end loop
```

---

- $I_{up} := \{i | r^{(i)} \alpha_i < B^{(i)}\}$  and  $I_{down} := \{j | A^{(j)} < r^{(j)} \alpha_j\}$  are indexes of  $\alpha_t$ 's who's values can still be moved up and down along the directions  $r^{(t)}$  and  $-r^{(t)}$  respectively

# Outline

- 1 Sparse Kernel Machines
- 2 Support Vector Regression
  - Primal Objective
- 3 Support Vector Classification**
  - Primal Objective
  - Dual Problem
  - Sequential Minimal Optimization
  - Termination Criteria
  - Directional Search**
  - Working Set Selection
  - Kernel Caching
- 4 Structural Risk Minimization

# Directional Search

- Suppose there is no constraint in the dual problem
- Let  $\alpha^{new} := \alpha + \lambda d$ , where  $d$  is a search direction
- Linear search:  $\lambda^* := \arg \max_{\lambda \geq 0} dual(\alpha + \lambda d)$  can be easily solved, yielding  $\lambda^* = \max(0, \frac{g^\top d}{d^\top K d})$  [Proof]
- When  $d$  points to an ascending direction (or, equivalently,  $g^\top d \geq 0$ ),  
$$\lambda^* = \frac{g^\top d}{d^\top K d}$$
  - This is true in SMO (to be explained in working set selection)

# Directional Search in SMO (1)

- Assume that the indexes of working set  $(i, j)$  is known (explained in working set selection), SMO picks the search direction  $\mathbf{d}^{(i,j)}$  that moves  $\boldsymbol{\alpha}$  along the  $i$ th and  $j$ th coordinates only (i.e.,  $d_t^{(i,j)} = 0$  for all  $t \neq i, j$ )
- In the presence of box constraint  $\mathbf{0} \leq \boldsymbol{\alpha} \leq \mathbf{C}$  and zero-sum constraint  $\mathbf{r}^\top \boldsymbol{\alpha} = 0$ , we need to ensure that  $\boldsymbol{\alpha}^{new} = \boldsymbol{\alpha} + \lambda \mathbf{d}$  is still feasible

# Directional Search in SMO (2)


- SMO deals with the zero-sum constraint by letting

$$d_t^{(i,j)} = \begin{cases} r^{(t)}, & \text{if } t = i \\ -r^{(t)}, & \text{if } t = j \\ 0, & \text{otherwise} \end{cases}$$

- So,  $\mathbf{r}^\top \boldsymbol{\alpha}^{new} = r^{(1)} \alpha_1 + \dots + r^{(i)} (\alpha_i + \lambda r^{(i)}) + \dots + r^{(j)} (\alpha_j - \lambda r^{(j)}) + \dots + r^{(N)} \alpha_N = \mathbf{r}^\top \boldsymbol{\alpha} + \lambda - \lambda = 0$
- Line search can be simplified:  $\lambda^* = \frac{\mathbf{g}^\top \mathbf{d}^{(i,j)}}{\mathbf{d}^{(i,j)\top} \tilde{\mathbf{K}} \mathbf{d}^{(i,j)}} = \frac{r^{(i)} g_i - r^{(j)} g_j}{\mathbf{K}_{i,i} + \mathbf{K}_{j,j} - 2\mathbf{K}_{i,j}}$

# Directional Search in SMO (3)

- SMO respects the box constraint by letting
$$\lambda^* = \min(B^{(i)} - r^{(i)}\alpha_i, r^{(j)}\alpha_j - A^{(j)}, \frac{r^{(i)}g_i - r^{(j)}g_j}{K_{i,i} + K_{j,j} - 2K_{i,j}})$$
- So,  $r^{(i)}(\alpha_i + \lambda r^{(i)}) = r^{(i)}\alpha_i + \lambda \leq B^{(i)}$  and
$$r^{(j)}(\alpha_j - \lambda r^{(j)}) = r^{(j)}\alpha_j - \lambda \geq A^{(j)}$$



# Sequential Minimal Optimization

---

**Algorithm 6.2** SMO with Maximum Violating Pair working set selection

---

1: $\forall k \in \{1 \dots n\}$	$\alpha_k \leftarrow 0$	▷ Initial coefficients
2: $\forall k \in \{1 \dots n\}$	$g_k \leftarrow 1$	▷ Initial gradient
3: loop		
4: $i \leftarrow \arg \max_i y_i g_i$	subject to $y_i \alpha_i < B_i$	
5: $j \leftarrow \arg \min_j y_j g_j$	subject to $A_j < y_j \alpha_j$	▷ Maximal violating pair
6:   if $y_i g_i \leq y_j g_j$	stop.	▷ Optimality criterion (11)
7: $\lambda \leftarrow \min \left\{ B_i - y_i \alpha_i, y_j \alpha_j - A_j, \frac{y_i g_i - y_j g_j}{K_{ii} + K_{jj} - 2K_{ij}} \right\}$		▷ Direction search
8: $\forall k \in \{1 \dots n\}$	$g_k \leftarrow g_k - \lambda y_k K_{ik} + \lambda y_k K_{jk}$	▷ Update gradient
9: $\alpha_i \leftarrow \alpha_i + y_i \lambda$	$\alpha_j \leftarrow \alpha_j - y_j \lambda$	▷ Update coefficients
10: end loop		

---



# Outline

- 1 Sparse Kernel Machines
- 2 Support Vector Regression
  - Primal Objective
- 3 Support Vector Classification**
  - Primal Objective
  - Dual Problem
  - Sequential Minimal Optimization
  - Termination Criteria
  - Directional Search
  - Working Set Selection**
  - Kernel Caching
- 4 Structural Risk Minimization

# Maximizing the Gain

- How to select the indexes  $(i,j)$  for the working set?

# Maximizing the Gain

- How to select the indexes  $(i, j)$  for the working set?
- Idea: pick the one that maximizes the gain:  
 $\mathbf{d}^{(i,j)*} := \arg \max_{\mathbf{d}^{(i,j)}} \max_{\lambda \geq 0} \text{dual}(\boldsymbol{\alpha} + \lambda \mathbf{d}^{(i,j)}) - \text{dual}(\boldsymbol{\alpha})$  subject to  
 $r^{(i)} \alpha_i + \lambda \leq B^{(i)}$  and  $A^{(j)} \leq r^{(j)} \alpha_j - \lambda$
- Unfortunately the search of the best direction  $\mathbf{d}^{(i,j)*}$  requires iterating over the  $N(N-1)$  possible pairs of indexes
  - The maximization of  $\lambda$  then amounts to performing a direction search
  - These repeated direction searches would virtually access all the kernel matrix during each SMO iteration
  - Not acceptable for a fast algorithm

# Maximal Violating Pair

- By the first order approximation, we have  $dual(\boldsymbol{\alpha} + \lambda \mathbf{d}^{(i,j)}) - dual(\boldsymbol{\alpha}) \approx \lambda \mathbf{g}^\top \mathbf{d}^{(i,j)}$  if  $\lambda$  is sufficiently small
- $\mathbf{d}^{(i,j)*} = \arg \max_{\mathbf{d}^{(i,j)}} \max_{0 \leq \lambda \leq \epsilon} \lambda \mathbf{g}^\top \mathbf{d}^{(i,j)}$
- We can assume that there is a direction  $\mathbf{d}^{(i,j)}$  such that  $\mathbf{g}^\top \mathbf{d}^{(i,j)} > 0$  because we would otherwise have reached the optimum
- $\mathbf{d}^{(i,j)*} = \arg \max_{\mathbf{d}^{(i,j)}} \mathbf{g}^\top \mathbf{d}^{(i,j)} = \arg \max_{i \in I_{up}, j \in I_{down}} (\mathbf{r}^{(i)} g_i - \mathbf{r}^{(j)} g_j)$
- Therefore,  $i := \max_{k \in I_{up}} \mathbf{r}^{(k)} g_k$  and  $j := \min_{k \in I_{down}} \mathbf{r}^{(k)} g_k$

# Sequential Minimal Optimization

---

**Algorithm 6.2** SMO with Maximum Violating Pair working set selection

---

1: $\forall k \in \{1 \dots n\} \quad \alpha_k \leftarrow 0$	▷ Initial coefficients
2: $\forall k \in \{1 \dots n\} \quad g_k \leftarrow 1$	▷ Initial gradient
3: loop	
4: $i \leftarrow \arg \max_i y_i g_i$ subject to $y_i \alpha_i < B_i$	
5: $j \leftarrow \arg \min_j y_j g_j$ subject to $A_j < y_j \alpha_j$	▷ Maximal violating pair
6:   if $y_i g_i \leq y_j g_j$ stop.	▷ Optimality criterion (11)
7: $\lambda \leftarrow \min \left\{ B_i - y_i \alpha_i, y_j \alpha_j - A_j, \frac{y_i g_i - y_j g_j}{K_{ii} + K_{jj} - 2K_{ij}} \right\}$	▷ Direction search
8: $\forall k \in \{1 \dots n\} \quad g_k \leftarrow g_k - \lambda y_k K_{ik} + \lambda y_k K_{jk}$	▷ Update gradient
9: $\alpha_i \leftarrow \alpha_i + y_i \lambda \quad \alpha_j \leftarrow \alpha_j - y_j \lambda$	▷ Update coefficients
10: end loop	

---

- The selected  $i$  and  $j$  are useful for **both** termination check and working set selection
- Compare the running time of SMO with that of CVX [Homework]

# Outline

- 1 Sparse Kernel Machines
- 2 Support Vector Regression
  - Primal Objective
- 3 Support Vector Classification**
  - Primal Objective
  - Dual Problem
  - Sequential Minimal Optimization
  - Termination Criteria
  - Directional Search
  - Working Set Selection
  - Kernel Caching
- 4 Structural Risk Minimization

# The Cost of SMO

---

**Algorithm 6.2** SMO with Maximum Violating Pair working set selection

---

1: $\forall k \in \{1 \dots n\} \quad \alpha_k \leftarrow 0$	▷ Initial coefficients
2: $\forall k \in \{1 \dots n\} \quad g_k \leftarrow 1$	▷ Initial gradient
3: loop	
4: $i \leftarrow \arg \max_i y_i g_i$ subject to $y_i \alpha_i < B_i$	
5: $j \leftarrow \arg \min_j y_j g_j$ subject to $A_j < y_j \alpha_j$	▷ Maximal violating pair
6:   if $y_i g_i \leq y_j g_j$ stop.	▷ Optimality criterion (11)
7: $\lambda \leftarrow \min \left\{ B_i - y_i \alpha_i, y_j \alpha_j - A_j, \frac{y_i g_i - y_j g_j}{K_{ii} + K_{jj} - 2K_{ij}} \right\}$	▷ Direction search
8: $\forall k \in \{1 \dots n\} \quad g_k \leftarrow g_k - \lambda y_k K_{ik} + \lambda y_k K_{jk}$	▷ Update gradient
9: $\alpha_i \leftarrow \alpha_i + y_i \lambda \quad \alpha_j \leftarrow \alpha_j - y_j \lambda$	▷ Update coefficients
10: end loop	

---

- Which line costs the most time?

# The Cost of SMO

---

**Algorithm 6.2** SMO with Maximum Violating Pair working set selection

---

1: $\forall k \in \{1 \dots n\} \quad \alpha_k \leftarrow 0$	▷ Initial coefficients
2: $\forall k \in \{1 \dots n\} \quad g_k \leftarrow 1$	▷ Initial gradient
3: loop	
4: $i \leftarrow \arg \max_i y_i g_i$ subject to $y_i \alpha_i < B_i$	
5: $j \leftarrow \arg \min_j y_j g_j$ subject to $A_j < y_j \alpha_j$	▷ Maximal violating pair
6:   if $y_i g_i \leq y_j g_j$ stop.	▷ Optimality criterion (11)
7: $\lambda \leftarrow \min \left\{ B_i - y_i \alpha_i, y_j \alpha_j - A_j, \frac{y_i g_i - y_j g_j}{K_{ii} + K_{jj} - 2K_{ij}} \right\}$	▷ Direction search
8: $\forall k \in \{1 \dots n\} \quad g_k \leftarrow g_k - \lambda y_k K_{ik} + \lambda y_k K_{jk}$	▷ Update gradient
9: $\alpha_i \leftarrow \alpha_i + y_i \lambda \quad \alpha_j \leftarrow \alpha_j - y_j \lambda$	▷ Update coefficients
10: end loop	

---

- Which line costs the most time? Line 8: updating of the gradient, which requires the computation of two full rows of the kernel matrix



# Computation of Kernel Values

- Computing the kernel values  $\mathbf{K}_{ij} = K(\mathbf{x}_i, \mathbf{x}_j)$  only appear as “constant factors” in the asymptotic complexity of solving the SVM problem
- But in practice it accounts for *more than half* the total running time!

# Observations

- Computing kernels is expensive: data can be high-dimensional, e.g., images have thousands of pixels, documents have thousands of words

# Observations

- Computing kernels is expensive: data can be high-dimensional, e.g., images have thousands of pixels, documents have thousands of words
- Computing the full kernel matrix is wasteful: the expression of the gradient  $g^{(i)} = 1 - r^{(i)} \sum_{j=1}^N r^{(j)} \alpha_j \mathbf{K}_{i,j}$  only depend on kernel values  $\mathbf{K}_{i,j}$  that involve at least one support vector (the other kernel values are multiplied by zero)

# Observations

- Computing kernels is expensive: data can be high-dimensional, e.g., images have thousands of pixels, documents have thousands of words
- Computing the full kernel matrix is wasteful: the expression of the gradient  $g^{(i)} = 1 - r^{(i)} \sum_{j=1}^N r^{(j)} \alpha_j \mathbf{K}_{i,j}$  only depend on kernel values  $\mathbf{K}_{i,j}$  that involve at least one support vector (the other kernel values are multiplied by zero)
- The kernel matrix may not fit in memory: large-scale problems are common today

# Kernel Caching

- In practice, SMO solvers implement a *cache* that keeps the most recently used kernel values
- A kernel value is computed on the fly when it is missed from cache
- Kernel cache hit rate becomes a major factor of the training time
- How much is the speedup? [Homework bonus]

# Outline

- 1 Sparse Kernel Machines
- 2 Support Vector Regression
  - Primal Objective
- 3 Support Vector Classification
  - Primal Objective
  - Dual Problem
  - Sequential Minimal Optimization
  - Termination Criteria
  - Directional Search
  - Working Set Selection
  - Kernel Caching
- 4 Structural Risk Minimization

# Structural Risk Minimization

TBA